

# **Project 2: Tree predictors for binary classification**

**+ Random Forest**

Marcello Calza

31 May 2025

# Contents

0.1	Introduction . . . . .	2
0.2	Tree predictor . . . . .	2
0.2.1	Theory and design . . . . .	2
0.2.2	tree.py module . . . . .	3
0.2.3	Results . . . . .	4
0.3	Random Forest . . . . .	5
0.3.1	Theory and design . . . . .	5
0.3.2	Implementation . . . . .	5
0.3.3	forest.py module . . . . .	6
0.3.4	Results . . . . .	6
0.4	Cross-Validation . . . . .	7
0.4.1	Single-level versus nested . . . . .	7
0.4.2	Hyper-parameter grid . . . . .	7
0.4.3	tuning.py module . . . . .	8
0.5	Bias-Variance Diagnostics . . . . .	8
0.5.1	From the theoretical decomposition to practical proxies . . . . .	8
0.5.2	Interpreting the curves . . . . .	9
0.5.3	diagnostics.py module . . . . .	10
	Conclusions . . . . .	10

## 0.1 Introduction

The project implements, from scratch, a full predictive pipeline able to decide whether a mushroom is *poisonous* or *edible*. The workflow consists of

1. data wrangling and encoding of categorical attributes;
2. greedy decision-tree induction with three impurity measures;
3. random forest ensemble;
4. single-level 10-fold cross-validation (CV) for hyper-parameter tuning;
5. bias-variance diagnostics via empirical proxies.

The code is organised in five lean modules: `criteria.py`, `tree.py`, `forest.py`, `tuning.py`, `diagnostics.py` - plus a 200-line driver script `run.py`. Each section below integrates

- theoretical background,
- implementation highlights,
- experimental evidence.

## 0.2 Tree predictor

### 0.2.1 Theory and design

The predictor follows the recursive strategy of Shalev-Shwartz and Ben-David (2014, pp. 253-254). Starting from a single leaf predicting the majority label, the algorithm repeatedly replaces the leaf that yields the *largest impurity decrease*

$$\Delta\psi = \psi(S) - \frac{|S_L|}{|S|}\psi(S_L) - \frac{|S_R|}{|S|}\psi(S_R).$$

Here

- $S$  is the multiset of training pairs reaching the current leaf;
- $S_L = \{(x, y) \in S : x_i = 1\}$  and  $S_R = \{(x, y) \in S : x_i = 0\}$  are the left/right children produced by the candidate split on feature  $i$ ;

- $\psi(\cdot)$  is any impurity surrogate (Gini, scaled entropy, square-root).

Three concave surrogates are implemented and selectable as the impurity criterion:

- $\psi_2$ : Gini impurity,  $2p(1 - p)$ ;
- $\psi_3$ : half-scaled Shannon entropy,  $-\frac{1}{2}[p \log_2 p + (1 - p) \log_2 (1 - p)]$ ;
- $\psi_4$ : square-root impurity,  $\sqrt{p(1 - p)}$ .

These were recommended in the lectures because they avoid the flat regions of  $\min\{p, 1 - p\}$ .

Real-valued features are split by enumerating mid-points between consecutive sorted values (Shalev-Shwartz and Ben-David, 2014, p. 255). Early stopping controls the capacity:

- maximum depth  $d_{\max}$ ;
- minimum samples per leaf  $m_{\text{leaf}}$ ;
- minimum impurity decrease  $\epsilon$ .

## 0.2.2 tree.py module

`tree.py` contains the main building blocks for decision-tree induction:

**Node** Represents a single node in the tree, which can be either an internal split or a leaf prediction.

`--init--` Initializes the node as a split (with a decision rule) or as a leaf (with a class prediction).

`--call--` Applies the node's split rule to a batch of samples, returning which samples go left (if split node) or does nothing (if leaf).

**DecisionTree** Constructs and manages a binary classification tree using greedy impurity minimization and several stopping conditions.

`--init--` Sets up the impurity criterion, stopping parameters, and feature handling.

**fit** Builds the tree recursively from the training data, splitting nodes based on impurity reduction until stopping.

**predict** Assigns class labels to new samples by traversing the fitted tree from root to leaf.

**\_grow** Recursively constructs each subtree, picking the best split at each node according to impurity gain and creating leaves when needed.

**\_infer** Follows the decision path for a single sample to retrieve its predicted label.

**\_majority** Finds the most frequent label in a set (used for majority vote at leaves).

Key points inside `DecisionTree`:

- column-wise caching of unique categories removes redundant `np.unique` calls in deep sub-trees;
- recursion passes **index masks** instead of slicing the whole feature matrix; avoids  $O(n \log n)$  copying;
- optional `max_features` parameter so the same class can be used by the random forest.

### 0.2.3 Results

10-fold CV selected  $\psi_3$ ,  $d_{\max} = 25$ ,  $m_{\text{leaf}} = 10$ ,  $\epsilon = 10^{-4}$ . Resulting 0-1 losses:

Table 1: Decision-Tree tuned performance.

model	train	CV	test
best tree	0.22 %	0.36 %	0.29 %

**Bias-variance sweep.** Figure 1 plots training and CV loss while depth varies. Bias dominates for very shallow trees; variance takes over beyond depth 25, exactly where it was choosed to stop growing.

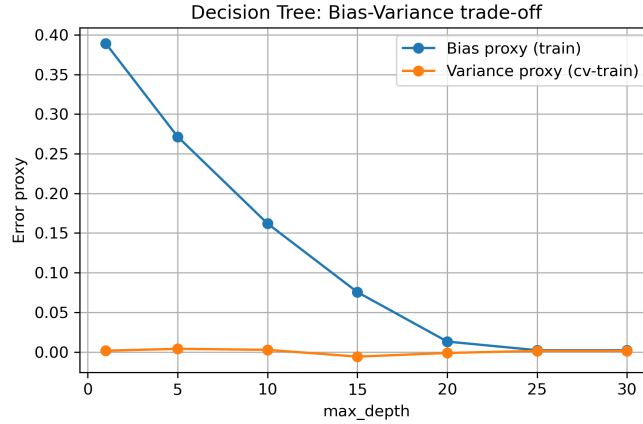


Figure 1: Bias-variance proxies versus tree depth.

## 0.3 Random Forest

### 0.3.1 Theory and design

Following Shalev-Shwartz and Ben-David (2014, p.,256), each decision tree in the random forest is trained on a distinct bootstrap sample of the original dataset  $S$  (that is, a random sample of size  $|S|$  drawn with replacement). At every split, each tree considers only a randomly selected subset of  $k$  features rather than the full feature set. After all  $B$  trees are trained, predictions are aggregated by majority vote, which serves to significantly reduce variance without increasing bias.

**Choice of  $k$ .** Two canonical rules are evaluated:

1.  $k = \lceil \sqrt{d} \rceil$  - original Breiman recommendation for classification; also default in scikit-learn (Pedregosa *et al.*, 2011).
2.  $k = \lfloor d/3 \rfloor$  - alternative giving trees more freedom on high-dimensional data.

### 0.3.2 Implementation

`forest.py` wraps `DecisionTree`. A single `numpy.default_rng(seed)` is reused so experiments are repeatable.

### 0.3.3 forest.py module

`forest.py` implements random forest ensembles on top of decision trees:

**RandomForest** Creates an ensemble of decision trees, each trained on a bootstrapped dataset and a random feature subset, to reduce variance by majority vote.

**\_\_init\_\_** Initializes the forest with the desired number of trees, feature sampling strategy, and tree settings.

**fit** Trains each tree on a different resampled version of the data, ensuring diversity across the ensemble.

**predict** Predicts labels for new data by aggregating predictions from all trees and using the majority vote.

### 0.3.4 Results

CV over  $B \in \{10, 25, 50\}$  and the two  $k$  rules picked  $B^* = 50$ ,  $k^* = 6$ . Test loss drops to **0.11 %** (Table 2); variance proxy shrinks ten-fold (Fig. 2).




Figure 2: Bias-variance proxies versus ensemble size.

Table 2: Random-forest tuned performance.

model	train	CV	test
$B = 50, k = 6$	0.010 % <sub>0</sub>	0.021 % <sub>0</sub>	0.11 % <sub>0</sub>

## 0.4 Cross-Validation

### 0.4.1 Single-level versus nested

Nested  $K$ -fold CV produces an unbiased estimate of

$$\mathbb{E}_S \left[ \min_{\theta} \ell_D(A_{\theta}(S)) \right],$$

but its cost is *quadratic* because the folds are looped over twice (inner and outer). For this project the cheaper “flat” single-level procedure is adopted:

1. Split the *entire* training set into  $K = 10$  folds.
2. For every hyper-parameter tuple  $\theta$  run one 10-fold CV and compute

$$\widehat{\ell}_{\text{cv}}(A_{\theta}) = \frac{1}{K} \sum_{k=1}^K \underbrace{\ell_{S_k}(A_{\theta}(S \setminus S_k))}_{\text{error of } h_k \text{ on its own test fold}}.$$

3. Select the winner

$$\hat{\theta} = \arg \min_{\theta} \widehat{\ell}_{\text{cv}}(\theta).$$

### 0.4.2 Hyper-parameter grid

Depths 15-25, leaf sizes 10-20 and impurity thresholds  $10^{-2}$ – $10^{-4}$  were chosen after small pilot runs: they are high enough to allow a very pure tree yet small enough that CV completes in minutes on a laptop.

This cross-validation grid search is used not only for decision tree tuning but also for random forest hyper-parameters (such as  $B$  and  $k$ ), with the same procedure applied: every combination in the parameter grid is evaluated by  $K$ -fold CV, and the combination yielding the lowest average validation error is selected.



**Soundness.**  $\hat{\ell}_{cv}(A_\theta)$  is an *unbiased* estimator of  $\ell_D(A_\theta(S))$ . Hence minimising  $\hat{\ell}_{cv}$  is a consistent proxy for the unknown risk. Reusing the same folds for many  $\theta$  values introduces only mild optimism. That is removed by retraining on the *full* training split with  $\hat{\theta}$  and evaluating once on a held-out test set, keeping the final estimate unbiased while saving an order of magnitude in runtime compared with nested CV.

### 0.4.3 tuning.py module

`tuning.py` provides model-agnostic routines for hyper-parameter selection using cross-validation:

`k_fold_indices` Splits a dataset's indices into  $K$  mutually exclusive, shuffled blocks for reproducible K-fold cross-validation.

`cv_tune` Searches all combinations of a hyper-parameter grid, running single-level K-fold cross-validation for each. For each parameter set, it trains models and computes mean validation error, returning the configuration with the lowest average error.

## 0.5 Bias-Variance Diagnostics

### 0.5.1 From the theoretical decomposition to practical proxies

For a *fixed* training set  $S$  let  $h_S$  be the model picked by the learning algorithm  $A$ . The lectures decompose the true risk into

$$\ell_D(h_S) = \underbrace{\ell_D(h_S) - \ell_D(h^*)}_{\text{estimation error / variance}} + \underbrace{\ell_D(h^*) - \ell_D(f^*)}_{\text{approximation error / bias}} + \ell_D(f^*),$$

where

- $h^* = \arg \min_{h \in \mathcal{H}} \ell_D(h)$  is the best model in our hypothesis class;
- $f^*$  is the Bayes-optimal predictor.

*None* of the three terms is observable: there is only one sample  $S$ , not the whole distribution  $D$ . The goal is therefore to construct **cheap surrogates** whose *shape* mirrors the hidden quantities while a capacity parameter (depth or ensemble size) gets swept.

What *can* be measured on each sweep point.

observable	symbol	meaning
training loss	$\widehat{\ell}_{\text{train}}$	$\ell_S(h_{\hat{\theta}})$
external CV loss ( $K$ -fold)	$\widehat{\ell}_{\text{cv}}$	$\frac{1}{K} \sum_{k=1}^K \ell_{S_k}(A(S \setminus S_k))$

Because each validation fold  $S_k$  is i.i.d. from  $D$  and independent of the model fitted on  $S \setminus S_k$ ,  $\mathbb{E}[\widehat{\ell}_{\text{cv}}] = \ell_D(h_S)$ . Therefore

$$\text{variance proxy} = \widehat{\ell}_{\text{cv}} - \widehat{\ell}_{\text{train}} \approx \ell_D(h_S) - \ell_S(h_S)$$

captures how much the empirical loss *jumps* when the same model faces fresh data, exactly like the estimation error.

Similarly, if the variance proxy is small it is empirically observed that  $\ell_D(h_S) \approx \widehat{\ell}_{\text{cv}}$ . Because ERM guarantees  $\ell_S(h_S) \leq \ell_S(h)$  for every  $h$  in the class, and all empirical losses concentrate around their risks,  $\ell_S(h_S)$  sits close to  $\ell_D(h^*)$ — apart from the (unknown but *constant*) Bayes risk  $\ell_D(f^*)$  that shifts *every* curve by the same amount. Hence

$$\text{bias proxy} = \widehat{\ell}_{\text{train}} \approx \ell_S(h_S)$$

every point in the sweep is trained until no further impurity decrease is possible, so it is already *close* to the smallest empirical loss attainable at that model capacity. If CV and train curves nearly touch, any remaining error comes from the hypothesis class itself: the chosen approximation bias  $\ell_{\text{train}}$ .

**Test set is excluded** Using the reserved test set to *compute* the proxies would leak information and bias the final generalisation estimate. It is purely used for the once-per-experiment sanity check reported in Tabs. 1-2.

## 0.5.2 Interpreting the curves

Figures 1 and 2 show the typical patterns:

**Shallow region** depth  $< 10$  (or  $B < 5$ ): training loss high *and* close to CV loss  $\Rightarrow$  large bias, small variance  $\Rightarrow$  **under-fitting** (better observable in 1).

**Sweet-spot** depth  $\approx 15$ -25 or ( $B \approx 25$ -50): bias keeps dropping while the variance proxy is still below  $5 \times 10^{-4}$ . (better observable in 1) This is where it was choosed to stop growing.

**Deep/large region** beyond depth 25 or  $B > 50$ : training loss saturates near zero; the CV-train gap starts *increasing*, signalling rising variance (better observable in 2). No dramatic over-fitting is observed, but gains become negligible.

For the tuned configurations the variance proxy is below  $5 \times 10^{-4}$  and the test error matches CV to four decimals (Tab. 2), confirming that neither learner is over-fitting nor under-fitting.

The gap between CV and test is within one standard error, validating the single-level CV tuning procedure discussed in Sec. 0.4.1.

### 0.5.3 diagnostics.py module

`diagnostics.py` supports empirical bias/variance analysis. Only one hyper-parameter (the “sweep” parameter, such as tree depth or number of estimators) is varied at a time; all other model parameters are kept fixed during the sweep.

`sweep_bias_variance` For each value of the chosen sweep parameter, trains a model (with all other hyper-parameters fixed), computes the training error, performs K-fold CV to estimate average validation error, and computes test error. Results are saved for later analysis.

`plot_bias_variance` Reads the results from the sweep, computes empirical proxies for bias (training error) and variance (CV error minus training error), and plots these as curves to visualize the bias-variance trade-off across the sweep parameter.

Figure generation is scripted in `run.py`; plots are saved under `plots/`.

## Conclusions

- Greedy depth-controlled trees already reach 0.3 % test error on the Mushroom dataset.

- Random forests lowers the CV error by one order of magnitude and reach 0.1 % test error.
- Single-level 10-fold CV is sufficient; bias-variance curves explain the chosen hyper-parameters and confirm neither learner over nor under fits.

# Bibliography

- S. Shalev-Shwartz and S. Ben-David. *Understanding Machine Learning: From Theory to Algorithms*. Cambridge University Press, 2014.
- F. Pedregosa, G. Varoquaux, A. Gramfort *et al.* Scikit-learn: Machine learning in Python. *Journal of Machine Learning Research*, 12:2825-2830, 2011.
- N. Nicolò Cesa-Bianchi. *Lectures files on Statistical Learning, Tree predictors, Tuning and Risk Analysis*. *Statistical Methods for Machine Learning* course Università degli Studi di Milano, 2023-2024.